

Lab 1. Puppet Language and Style



We will cover the following recipes in this lab:

- Adding a resource to a node
- Using `facter` to describe a node
- Using Puppet facts
- Installing a package before starting a service
- Installing, configuring, and starting a service
- Using community Puppet style
- Installing Puppet
- Creating a manifest
- Checking your manifests with Puppet-lint
- Making modules
- Using standard naming conventions

Introduction

In this lab, we'll start with the basics of the Puppet syntax and show you how some of the syntactic sugar in Puppet is used. We'll then move on to how Puppet deals with dependencies and how to make Puppet do the work for you.

We'll look at how to organize and structure your code into modules following community conventions so that other people will find it easy to read and maintain your code. We will also see some powerful features of the Puppet language, which will let you write concise yet expressive manifests.

Adding a resource to a node

This recipe will introduce the language and show you the basics of writing Puppet code. Puppet code files are called **manifests**; manifests declare resources. A resource in Puppet may be a type, class, or node. A **type** is something like a file or package or anything that has a type declared in the language. The current list of standard types is available on the puppetlabs website at <https://puppet.com/docs/puppet/latest/type.html>. I find myself referencing this site very often. You may define your own types, either using a mechanism, similar to a subroutine, named `defined types`, extending the language using a custom type. Types are the heart of the language; they describe the things that make up a **node** (node is the word Puppet uses for client computers/devices). Puppet uses **resources** to describe the state of a node; for example, we will declare the following package resource for a node using a site manifest: `site.pp`.

How to do it...

Create a `site.pp` file and place the following code in it:

```
node default {
  package { ['httpd']:
    ensure => 'installed'
  }
}
```

How it works...

This manifest will ensure that any node on which this manifest is applied will install a package called `httpd`. The `default` keyword is a wildcard to Puppet; it applies anything within the `node default` definition to any node. When Puppet applies the manifest to a node, it uses a **Resource Abstraction Layer (RAL)** to translate the package type into the package management system of the target node. What this means is that we can use the same manifest to install the `httpd` package on any system where Puppet has a Provider for the package type. Providers are the

pieces of code that do the real work of applying a manifest. When the previous code is applied to a node running on a YUM-based distribution, the YUM provider will be used to install the `httpd` RPM packages. When the same code is applied to a node running on an APT-based distribution, the APT provider will be used to install the `httpd` DEB package (which may not exist, as most Debian-based systems call this package `apache2` ; we'll deal with this sort of naming problem later).

Using `facter` to describe a node

Facter is a separate utility upon which Puppet depends. It is the system used by Puppet to gather information about the target system (node); `facter` calls the nuggets of information facts. You may run `facter` from the command line to obtain real-time information from the system.

How to do it...

We'll compare the output of `facter` with that of system utilities:

1. Use `facter` to find the current `uptime` of the system, the `uptime` fact:

```
t@fenago ~$ facter uptime 0:12 hours
```

2. Compare this with the output of the Linux `uptime` command:

```
t@fenago ~$ uptime
01:18:52 up 12 min, 1 user, load average: 0.00, 0.00, 0.00
```

How it works...

When `facter` is installed (as a dependency for Puppet), several fact definitions are installed by default. You can reference each of these facts by name from the command line.

There's more...

Running `facter` without any arguments causes `facter` to print all the facts known about the system. We will see in later chapters that `facter` can be extended with your own custom facts. All facts are available for you to use as variables; variables are discussed in the next section.

Variables

Variables in Puppet are marked with a `$` character. Variables are immutable; once assigned a value, they cannot be changed. When using variables within a manifest, it is advisable to enclose the variable within braces, such as `{myvariable}`, instead of `$myvariable`. All of the facts from `facter` can be referenced as top-scope variables (we will discuss scope in the next section). For example, the **Fully Qualified Domain Name (FQDN)** of the node may be referenced by `{::fqdn}`. Variables can only contain alphabetic characters, numerals, and the underscore character, `_`. As a matter of style, variables should start with an alphabetic character. Never use dashes in variable names.

Scope

In the variable example explained in the previous paragraph, the FQDN was referred to as `{::fqdn}` rather than `{fqdn}`; the double colons are how Puppet differentiates scope. The highest level scope, top-scope, or global is referred to by two colons, as in `::`, at the beginning of a variable identifier. To reduce namespace collisions, always use fully scoped variable identifiers in your manifests. A Unix user can think of top-scope variables such as the `/` (root) level. You can refer to variables using the double colon syntax, similar to how you would refer to a directory by

its full path. A developer can think of top-scope variables as global variables; however, unlike global variables, you must always refer to them with the double colon notation to guarantee that a local variable isn't obscuring the top-scope variable. In Puppet5, it is advisable to use the `$facts` fact, so the previous would be `${facts['fqdn']}`. When referring to a variable, the braces (`{}`) are optional outside of a string, as shown in the following example:

```
$fqdn_ = $facts['fqdn']
notify {"::fqdn is ${::fqdn}": }
notify {"fqdn_ is ${fqdn_}": }
notify {"facts['fqdn'] is ${facts['fqdn']}": }
```

This produces the following output:

```
t@mylaptop ~ $ puppet apply fqdn.pp
Notice: Compiled catalog for mylaptop.example.com in environment production in 0.01
seconds
Notice: ::fqdn is mylaptop.example.com
Notice: /Stage[main]/Main/Notify[::fqdn is mylaptop.example.com]/message: defined
'message' as '::fqdn is mylaptop.example.com'
Notice: fqdn_ is mylaptop.example.com
Notice: /Stage[main]/Main/Notify[fqdn_ is mylaptop.example.com]/message: defined
'message' as 'fqdn_ is mylaptop.example.com'
Notice: facts['fqdn'] is mylaptop.example.com
Notice: /Stage[main]/Main/Notify[facts['fqdn'] is mylaptop.example.com]/message:
defined 'message' as 'facts['fqdn'] is mylaptop.example.com'
Notice: Applied catalog in 0.02 seconds
```

Note

`$fqdn_` is used to avoid a namespace collision with the top-scope `::fqdn`.

Using puppet facts

As we'll see in subsequent chapters, `facter` may be extended with custom facts written in Ruby. By default, custom facts are not loaded when you run `facter`.

How to do it...

To pull in the custom facts, you need to specify the `-p` option to `facter`, as shown here:

```
t@fenago:~$ facter puppetversion

t@fenago:~$ facter -p puppetversion
5.5.2
```

Although still valid, the `facter -p` syntax is now deprecated in favor of using the Puppet face, `facts`. Puppet faces are the various sub-applications supported by the `Puppet` command. To see the available faces, run `Puppet help`, as shown here:

```
t@fenago:~$ puppet help

Usage: puppet <subcommand> [options] <action> [options]

Available subcommands:
```

| | |
|-----------------------------|---|
| agent | The puppet agent daemon |
| apply | Apply Puppet manifests locally |
| ca | Local Puppet Certificate Authority management. (Deprecated) |
| catalog | Compile, save, view, and convert catalogs. |
| cert | Manage certificates and requests |
| certificate | Provide access to the CA for certificate management. |
| certificate_request | Manage certificate requests. (Deprecated) |
| certificate_revocation_list | Manage the list of revoked certificates. (Deprecated) |
| config | Interact with Puppet's settings. |
| describe | Display help about resource types |
| device | Manage remote network devices |
| doc | Generate Puppet references |
| epp | Interact directly with the EPP template parser/renderer. |
| facts | Retrieve and store facts. |
| filebucket | Store and retrieve files in a filebucket |
| generate | Generates Puppet code from Ruby definitions. |
| help | Display Puppet help. |
| key | Create, save, and remove certificate keys. (Deprecated) |
| lookup | Interactive Hiera lookup |
| man | Display Puppet manual pages. |
| master | The puppet master daemon |
| module | Creates, installs and searches for modules on the Puppet Forge. |
| node | View and manage node definitions. |
| parser | Interact directly with the parser. |
| plugin | Interact with the Puppet plugin system. |
| report | Create, display, and submit reports. |
| resource | The resource abstraction layer shell |
| status | View puppet server status. (Deprecated) |

One difference between `facter` and `Puppet facts` is that you may request a single fact from `facter`, whereas `Puppet facts` will return all the facts for a node at once as a JSON object (you may request other formats with the `--render-as` option).

Installing a package before starting a service

To show how ordering works, we'll create a manifest that installs `httpd` and then ensures the `httpd` package service is running.

How to do it...

We'll create a manifest to install and start our service:

1. Start by creating a manifest that defines `service` :

```
service {'httpd':
  ensure => running,
  require => Package['httpd'],
}
```

2. The `service` definition references a `package` resource named `httpd`; we now need to define that resource:

```
package {'httpd':  
  ensure => 'installed',  
}
```

How it works...

In this example, the package will be installed before the service is started. Using `require` within the definition of the `httpd` service ensures that the package is installed first, regardless of the order within the manifest file.

Capitalization is important in Puppet. In our previous example, we created a package named `httpd`. If we wanted to refer to this package later, we would capitalize its type (package) as follows:

```
Package['httpd']
```

To refer to a class- for example, the `something::somewhere` class, which has already been included/defined in your manifest-you can reference it with the full path as follows:

```
Class['something::somewhere']
```

Let's say you have defined the following type:

```
example::thing {'one':}
```

The preceding line may be referenced later, as follows:

```
Example::Thing['one']
```

Knowing how to reference previously defined resources is necessary for the next section on metaparameters and ordering.

Learning metaparameters and ordering

All the manifests that will be used to define a node are compiled into a catalog. A catalog is the code that will be applied to configure a node. It is important to remember that manifests are not applied to nodes sequentially. There is no inherent order to the application of manifests. With this in mind, in the previous `httpd` example, what if we wanted to ensure that the `httpd` process started after the `httpd` package was installed?

We couldn't rely on the `httpd` service coming after the `httpd` package in the manifests. What we would have to do is use metaparameters to tell Puppet the order in which we want resources applied to the node. Metaparameters are parameters that can be applied to any resource and are not specific to any one resource type. They are used for catalog compilation and as hints to Puppet, but not to define anything about the resource to which they are attached.

When dealing with ordering, there are four metaparameters used:

- `before`
- `require`
- `notify`
- `subscribe`

The `before` and `require` metaparameters specify a direct ordering; `notify` implies `before` and `subscribe` implies `require`. The `notify` metaparameter is only applicable to services; what `notify` does is tell a service to restart after the notifying resource has been applied to the node (this is most often a package or file resource). In the case of files, once the file is created on the node, a `notify` parameter will restart any services

mentioned. The `subscribe` metaparameter has the same effect but is defined on the service; the service will `subscribe` to the file.

Trifecta

The relationship between package and service previously mentioned is an important and powerful paradigm of Puppet. Adding one more resource-type file into the fold creates what puppeteers refer to as the trifecta. Almost all system administration tasks revolve around these three resource types. As a system administrator, you install a package, configure the package with files, and then start the service:




Diagram of the trifecta (files require package for directory; service requires files and package)

Idempotency

A key concept of Puppet is that the state of the system when a manifest is applied to a node cannot affect the outcome of the Puppet run. In other words, at the end of the Puppet run (if the run was successful), the system will be in a known state and any further application of the manifest will result in a system that is in the same state. This property of Puppet is known as idempotency. Idempotency is the property that, no matter how many times you do something, remains in the same state as the first time you did it. For instance, if you had a light switch and you gave the instruction to turn it on, the light would turn on. If you gave the instruction again, the light would remain on.

Installing, configuring, and starting a service

There are many examples of this pattern online. In our simple example, we will create an Apache configuration file under `/etc/httpd/conf.d/fenago.conf`. The `/etc/httpd/conf.d` directory will not exist until the `httpd` package is installed. After this file is created, we would want `httpd` to restart to notice the change; we can achieve this with a `notify` parameter.

How to do it...

We will need the same definitions as our last example; we need the package and service installed. We now need two more things. We need the configuration file and index page (`index.html`) created. For this, we follow these steps:

1. As in the previous example, we ensure the service is running and specify that the service requires the `httpd` package:

```
service {'httpd':  
  ensure => running,  
  require => Package['httpd'],  
}
```

2. We then define `package` as follows:

```
package {'httpd':  
  ensure => installed,  
}
```

Note: Apache2 cannot start because port 80 is already in use. We will change the port first.

Change Apache Port

Let's change the default Apache port. Open `/etc/apache2/ports.conf` in vscode and update `80` with port `81`. Save and close the file.

3. Now, we create the `/etc/httpd/conf.d/fenago.conf` configuration file; the `/etc/httpd/conf.d` directory will not exist until the `httpd` package is installed. We'll use `@heredoc` syntax here to make the code a little more readable, assigning the `fenago.conf` contents to the `$fenago` variable. The `require` metaparameter tells Puppet that this file requires the `httpd` package to be installed before it is created:

```
$fenago = @(FENAGO)  
<VirtualHost *:81>  
  Servername fenago  
  DocumentRoot /var/www/fenago  
</VirtualHost>  
| FENAGO  
file {'/etc/httpd/conf.d/fenago.conf':  
  content => $fenago,  
  require => Package['httpd'],  
  notify => Service['httpd'],  
}
```

4. We then go on to create an `index.html` file for our virtual host in `/var/www/fenago`. Again, we'll use `@heredoc` syntax to make this more readable. This directory won't exist yet, so we need to create this as well, using the following code:

```
$index = @(INDEX)  
<html>  
  <body>  
    <h1>Hello World!</h1>  
  </body>  
</html>  
| INDEX  
file {'/var/www/fenago':  
  ensure => directory,  
  require => Package['httpd'],  
}
```

```
}
file {'/var/www/fenago/index.html':
  content => $index,
  require => File['/var/www/fenago'],
}
```

How it works...

The `require` attribute to the file resources tells Puppet that we need the `/var/www/fenago` directory created before we can create the `index.html` file. The important concept to remember is that we cannot assume anything about the target system (node). We need to define everything on which the target depends. Anytime you create a file in a manifest, you have to ensure that the directory containing that file exists. Anytime you specify that a service should be running, you have to ensure that the package providing that service is installed.

In this example, using metaparameters, we can be confident that no matter what state the node is in before running Puppet, after Puppet runs, the following will be true:

- `httpd` will be running
- The `VirtualHost` configuration file will exist
- `httpd` will restart and be aware of the `VirtualHost` file
- The `DocumentRoot` directory will exist
- An `index.html` file will exist in the `DocumentRoot` directory

Using community Puppet style

If other people need to read or maintain your manifests, or if you want to share code with the community, it's a good idea to follow the existing style conventions as closely as possible. These govern such aspects of your code as layout, spacing, quoting, alignment, and variable references, and the official puppetlabs recommendations on style are available at https://puppet.com/docs/puppet/latest/style_guide.html.

How to do it...

In this section, I'll show you a few of the more important examples and how to make sure that your code is style compliant.

Indentation

Indent your manifests using two spaces (not tabs), as follows:

```
service {'httpd':
  ensure => 'running',
}
```

Quoting

Always quote your resource names, as follows:

```
package { 'exim4': }
```

We cannot do this though:

```
package { exim4: }
```

Use single quotes for all strings, except when:

- The string contains variable references, such as `${::fqdn}`
- The string contains character escape sequences, such as `\n`

Consider the following code:

```
file { '/etc/motd':
  content => "Welcome to ${::fqdn}\n"
}
```

Puppet doesn't process variable references or escape sequences unless they're inside double quotes.

Always quote parameter values that are not reserved words in Puppet. For example, the following values are not reserved words:

```
name  => 'Nucky Thompson',
mode  => '0700',
owner => 'deploy',
```

However, these values are reserved words and therefore not quoted:

```
ensure => installed,
enable => true,
ensure => running,
```

false

There is only one thing in Puppet that is `false`, that is, the word `false` without any quotes. The string `false` evaluates to true and the string `true` also evaluates to true. Actually, everything besides the literal false evaluates to true (when treated as a Boolean):

```
if "false" {
  notify { 'True': }
}
if 'false' {
  notify { 'Also true': }
}
if false {
  notify { 'Not true': }
}
```

When this code is run through `puppet apply`, the first two notifies are triggered. The final notify is not triggered; it is the only one that evaluates to false.

Variables

Always include curly braces `{}` around variable names when referring to them in strings, for example, as follows:

```
source => "puppet:///modules/webserver/${brand}.conf",
```

Otherwise, the `Puppet` parser has to guess which characters should be a part of the variable name and which belong to the surrounding string. Curly braces make it explicit.

Parameters

Always end lines that declare parameters with a comma, even if it is the last parameter:

```

service { 'memcached':
  ensure => running,
  enable => true,
}

```

This is allowed by Puppet, and makes it easier if you want to add parameters later, or reorder the existing parameters.

When declaring a resource with a single parameter, make the declaration all on one line and with no trailing comma, as shown in the following snippet:

```

package { 'puppet':
  ensure => installed
}

```

Where there is more than one parameter, give each parameter its own line:

```

package { 'rake':
  ensure    => installed,
  provider  => gem,
  require   => Package['rubygems'],
}

```

To make the code easier to read, line up the parameter arrows in line with the longest parameter, as follows:

```

file { "/var/www/${app}/shared/config/rvmrc":
  owner    => 'deploy',
  group    => 'deploy',
  content  => template('rails/rvmrc.erb'),
  require  => File["/var/www/${app}/shared/config"],
}

```

The arrows should be aligned per resource, but not across the whole file, otherwise it may be difficult for you to cut and paste code from one file to another.

Symlinks

When declaring file resources that are symlinks, use the `ensure => link` link and set the target attribute, as follows:

```

file { '/etc/php5/cli/php.ini':
  ensure => link,
  target => '/etc/php.ini',
}

```

Installing Puppet

You may install Puppet locally on your machine or create a virtual machine and install Puppet on that machine.

How to do it...

For YUM-based systems, use <https://yum.puppetlabs.com/puppet5>, and for APT-based systems, use <https://apt.puppetlabs.com/>. Use the `puppet5-release-[version].deb` package to install Puppet 5. After installing the YUM release RPM or the APT release source package, install the puppet-agent package. The puppet-agent package installs all the files necessary to support Puppet in agent mode on a node.

You may also use `gem` to install Puppet:

1. To use `gem`, we need the `rubygems` package as follows:

```
t@fenago:~$ sudo yum install rubygems
Resolving Dependencies
--> Running transaction check
--> Package rubygems.noarch 0:2.0.14.1-30.el7 will be installed
...
Installed:
  rubygems.noarch 0:2.0.14.1-30.el7
```

2. Now, use `gem` to install Puppet:

```
t@fenago:~$ gem install puppet

Fetching: facter-2.5.1.gem (100%)
Successfully installed facter-2.5.1
Fetching: hiera-3.4.2.gem (100%)
Successfully installed hiera-3.4.2
Fetching: fast_gettext-1.1.0.gem (100%)
Successfully installed fast_gettext-1.1.0
Fetching: locale-2.1.2.gem (100%)
Successfully installed locale-2.1.2
Fetching: text-1.3.1.gem (100%)
Successfully installed text-1.3.1
Fetching: gettext-3.2.6.gem (100%)
Successfully installed gettext-3.2.6
Fetching: gettext-setup-0.29.gem (100%)
Successfully installed gettext-setup-0.29
Fetching: puppet-5.0.0.gem (100%)
Successfully installed puppet-5.5.2
Parsing documentation for facter-2.5.1
Installing ri documentation for facter-2.5.1
Parsing documentation for hiera-3.4.2
Installing ri documentation for hiera-3.4.2
Parsing documentation for fast_gettext-1.1.0
Installing ri documentation for fast_gettext-1.1.0
Parsing documentation for locale-2.1.2
Installing ri documentation for locale-2.1.2
Parsing documentation for text-1.3.1
Installing ri documentation for text-1.3.1
Parsing documentation for gettext-3.2.6
Installing ri documentation for gettext-3.2.6
Parsing documentation for gettext-setup-0.29
Installing ri documentation for gettext-setup-0.29
Parsing documentation for puppet-5.5.2
Installing ri documentation for puppet-5.5.2
8 gems installed
```

For the examples in this course, I suggest using the puppet-agent package installation. The package installation method of Puppet uses the **AIO (All-In-One)** mentality. The puppet-agent package installs all the necessary support

files for Puppet and does not rely on system libraries and applications. The most important dependency is Ruby: the AIO `puppet-agent` package installs a Puppet-specific Ruby that has been tested against the version of Puppet to which it belongs.

Creating a manifest

If you already have some Puppet code (known as a Puppet manifest), you can skip this section and go on to the next. If not, we'll see how to create and apply a simple manifest.

How to do it...

To create and apply a simple manifest, follow these steps:

1. With Puppet installed in the previous section, we can create a directory to contain our Puppet code:

```
t@fenago:~$ mkdir -p .puppet/manifests
t@fenago:~$ cd .puppet/manifests
t@fenago:manifests$
```

2. Within your manifests directory, create the `site.pp` file with the following content:

```
node default {
  file { ['/tmp/hello':
    content => "Hello, world!\n",
  ]
}
```

3. Test your manifest with the `puppet apply` command. This will tell Puppet to read the manifest, compare it to the state of the machine, and make any necessary changes to that state:

```
t@fenago:manifests$ puppet apply site.pp
Notice: Compiled catalog for fenago.example.com in environment production in 0.05
seconds
Notice: /Stage[main]/Main/Node[default]/File[/tmp/hello]/ensure: defined content as
'{md5}746308829575e17c3331bbcb00c0898b'
Notice: Applied catalog in 0.07 seconds
```

4. To see if Puppet did what we expected (created the `/tmp/hello` file with the Hello, world! content), run the following command:

```
t@fenago:manifests$ cat /tmp/hello
Hello, world!
```

Note that creating the file in `/tmp` did not require special permissions. We did not run Puppet via `sudo`. Puppet need not be run through `sudo`; there are cases where running via an unprivileged user can be useful.

There's more...

When several people are working on a code base, it's easy for style inconsistencies to creep in. Fortunately, there's a tool available that can automatically check your code for compliance with the style guide: `puppet-lint`. We'll see how to use this in the next section.

Checking your manifests with puppet-lint

The Puppet official style guide outlines a number of style conventions for Puppet code, some of which we've touched on in the preceding section. For example, according to the style guide, manifests:

- Must use two-space soft tabs
- Must not use literal tab characters
- Must not contain trailing white space
- Should not exceed an 80-character line width
- Should align parameter arrows (`=>`) within blocks

Following the style guide will make sure that your Puppet code is easy to read and maintain, and if you're planning to release your code to the public, style compliance is essential.

The `puppet-lint` tool will automatically check your code against the style guide. The next section explains how to use it.

Getting ready

Here's what you need to do to install `puppet-lint`:

1. We'll install Puppet-lint using the gem provider because the gem version is much more up to date than the APT or RPM packages available. Create a `puppet-lint.pp` manifest as shown in the following code snippet:

```
package {'puppet-lint':  
  ensure => 'installed',  
  provider => 'gem'  
}
```

2. Run `puppet apply` on the `puppet-lint.pp` manifest, as shown in the following command:

```
t@fenago:manifests$ puppet apply puppet-lint.pp  
Notice: Compiled catalog for fenago.example.com in environment production in 1.04  
seconds  
Notice: /Stage[main]/Main/Package[puppet-lint]/ensure: created  
Notice: Applied catalog in 0.93 seconds
```

How to do it...

Follow these steps to use Puppet-lint:

1. Choose a Puppet manifest file that you want to check with `Puppet-lint`, and run the following command:

```
t@fenago:manifests$ puppet-lint puppet-lint.pp  
WARNING: indentation of => is not properly aligned (expected in column 12, but found  
it in column 10) on line 2  
ERROR: trailing whitespace found on line 4
```

2. As you can see, `Puppet-lint` found a number of problems with the manifest file. Correct the errors, save the file, and rerun `Puppet-lint` to check that all is well. If successful, you'll see no output:

```
t@fenago:manifests$ puppet-lint puppet-lint.pp
t@fenago:manifests$
```

There's more...

Should you follow Puppet style guide and, by extension, keep your code lint-clean? It's up to you, but here are a couple of things to think about:

- It makes sense to use some style conventions, especially when you're working collaboratively on code. Unless you and your colleagues can agree on standards for whitespace, tabs, quoting, alignment, and so on, your code will be messy and difficult to read or maintain.
- If you're choosing a set of style conventions to follow, the logical choice would be those issued by Puppet and adopted by the community for use in public modules.

Having said that, it's possible to tell `Puppet-lint` to ignore certain checks if you've chosen not to adopt them in your code base. For example, if you don't want `puppet-lint` to warn you about code lines exceeding 80 characters, you can run `puppet-lint` with the following option:

```
puppet-lint --no-80chars-check
```

Most developers have terminals with more than 80 characters now; the check for 80 characters is generally disabled in favor of a new 140-character limit. You may disable the 140 character check with the following:

```
puppet-lint --no-140chars-check
```

Run `puppet-lint --help` to see the complete list of check configuration commands.

See also

- [You can find out more about] `Puppet-lint` [at] <https://github.com/rodjek/puppet-lint>.
- The [Automatic syntax checking with Git hooks] recipe in [Lab 2], [Puppet Infrastructure]
- The [Testing your manifests with rspec-puppet] recipe in [Lab 9], [External Tools and the Puppet Ecosystem]

Making modules

One of the most important things you can do to make your Puppet manifests clearer and more maintainable is to organize them into modules.

Modules are self-contained bundles of Puppet code that include all the files necessary to implement a thing. Modules may contain flat files, templates, Puppet manifests, custom fact declarations, Augeas lenses, and custom Puppet types and providers.

Separating things into modules makes it easier to reuse and share code; it's also the most logical way to organize your manifests. In this example, we'll create a module to manage `memcached`, a memory caching system commonly used with web applications.

How to do it...

Following are the steps to create an example module:

1. We will use Puppet's module subcommand to create the directory structure for our new module, in our home directory (`/home/vagrant`):

```
[t@fenago ~]$ puppet module generate thomas-memcached
```

2. We need to create a `metadata.json` file for this module. Please answer the following questions; if the question is not applicable to this module, feel free to leave it blank:

```
Puppet uses Semantic Versioning (semver.org) to version modules.What version is this
module? [0.1.0]-->Who wrote this module? [thomas]-->What license does this module code
fall under? [Apache-2.0]-->How would you describe this module in a single sentence?-->
A module to install memcachedWhere is this module's source code repository?-->
github.com/uphillian/thomas-memcachedWhere can others go to learn more about this
module? [https://github.com/uphillian/thomas-memcached]-->Where can others go to file
issues about this module? [https://github.com/uphillian/thomas-memcached/issues]-->
{"name": "thomas-memcached","version": "0.1.0","author": "thomas","summary": "A module
to install memcached","license": "Apache-2.0","source": "github.com/uphillian/thomas-
memcached","project_page":
"https://github.com/uphillian/thosasmemcached","issues_url":
"https://github.com/uphillian/thomas-memcached/issues","dependencies": [{"name":
"puppetlabs-stdlib","version_requirement": ">= 1.0.0"}]"data_provider": null}-----
-----About to generate this metadata; continue? [n/Y]-->
yNotice: Generating module at /home/vagrant/memcached...Notice: Populating
templates...Finished; module generated in
memcached.memcached/specmemcached/spec/spec_helper.rbmemcached/spec/classesmemcached/spe
```

This command creates the module directory and creates some empty files as starting points.

3. Now, edit `memcached/manifests/init.pp` and change the class definition at the end of the file to the following. Note that the `puppet module` created many lines of comments; in a `production` module, you would want to edit those default comments:

```
class memcached {
  package { ['memcached': ensure => installed, ]
  file { ['/etc/memcached.conf':
    source => 'puppet:///modules/memcached/memcached.conf',
    owner  => 'root',
    group  => 'root',
    mode   => '0644',
    require => Package['memcached'],
  }
  service { ['memcached':
    ensure => running,
    enable => true,
    require => [Package['memcached'],
               File['/etc/memcached.conf']],
  }
}
```

4. Create the `modules/thomas-memcached/files` directory and then create a file named `memcached.conf` with the following contents:

```
[t@fenago memcached]$ mkdir files
[t@fenago memcached]$ echo "-m 64 -p 11211 -u nobody -l 127.0.0.1" >
files/memcached.conf
```

5. We would like this module to install `memcached`. We'll need to run `Puppet` with root privileges, and we'll use `sudo` for that. We'll need `Puppet` to be able to find the module in our home directory; we can specify this on the command line when we run `Puppet`, as shown in the following code snippet:

```
t@fenago:memcached$ sudo /opt/puppetlabs/bin/puppet apply --modulepath=/home/vagrant -
e 'include memcached'Warning: ModuleLoader: module 'memcached' has unresolved
dependencies - it will only see those that are resolved. Use 'puppet module list --
tree' to see information about modules (file & line not available)Notice:
Compiled catalog for fenago.strangled.net in environment production in 0.46
secondsNotice: /Stage[main]/Memcached/Package[memcached]/ensure: createdNotice:
/Stage[main]/Memcached/File[/etc/memcached.conf]/ensure: defined content as
'{md5}febccf4a987759cf4f1558cc625fbea9'Notice:
/Stage[main]/Memcached/Service[memcached]/ensure: ensure changed 'stopped' to
'running'Notice: Applied catalog in 6.99 seconds
```

6. We can verify that `memcached` is running using `systemctl` or `puppet resource`:

```
t@fenago:memcached$ sudo /opt/puppetlabs/bin/puppet resource service memcached
service { 'memcached':
  ensure => 'running',
  enable => 'true',
}
t@fenago:memcached$ sudo systemctl status memcached
memcached.service - Memcached
   Loaded: loaded (/usr/lib/systemd/system/memcached.service; enabled; vendor preset:
disabled)
   Active: active (running) since Thu 2017-12-28 05:17:41 UTC; 3min 28s ago
 Main PID: 4057 (memcached)
    CGroup: /system.slice/memcached.service
            └─4057 /usr/bin/memcached -u memcached -p 11211 -m 64 -c 1024

Dec 28 05:17:41 fenago systemd[1]: Started Memcached.
Dec 28 05:17:41 fenago systemd[1]: Starting Memcached...
Note that /opt/puppetlabs/bin/puppet may not be in root's path, use the full path or
add the path to a file in /etc/profile.d.
```

How it works...

When we created the module using Puppet's `module generate` command, we used the name `thomas-memcached`. The name before the hyphen is your username or your username on Puppet forge (an online repository of modules). Modules have a specific directory structure. Not all of these directories need to be present, but if they are, this is how they should be organized:

```
modules/
└─MODULE_NAME/ never use a dash (-) in a module name
   └─examples/ example usage of the module
      └─files/ flat files used by the module
         └─lib/
            └─factor/ define new facts for factor
               └─puppet/
                  └─parser/
```



```
└functions/ define a new puppet function, like sort()
└provider/ define a provider for a new or existing type
└util/ define helper functions (in ruby)
└type/ define a new type in puppet
└manifests/
└init.pp class MODULE_NAME { }
└spec/ rSpec tests
└templates/ EPP or ERB template files used by the module
```

All manifest files (those containing Puppet code) live in the manifests directory. In our example, the memcached class is defined in the `manifests/init.pp` file, which will be imported automatically.

Inside the memcached class, we refer to the `memcached.conf` file:

```
file { ['/etc/memcached.conf':
  source => 'puppet:///modules/memcached/memcached.conf',
}
```

The preceding `source` parameter tells Puppet to look for the file in:

```
MODULEPATH/(/home/vagrant/)
└memcached/
└files/
└memcached.conf
```

There's more...

Learn to love modules because they'll make your Puppet life a lot easier. They're not complicated, however; practice and experience will help you judge when things should be grouped into modules, and how best to arrange your module structure. Modules can hold more than manifests and files, as we'll see in the next two sections.

Templates

If you need to use a template as a part of the module, place it in the module's templates directory and refer to it as follows:

```
file { ['/etc/memcached.conf':
  content => epp('memcached/memcached.conf.epp'),
}
```

Puppet will look for the file in:

```
MODULEPATH/memcached/templates/memcached.conf.epp
```

Facts, functions, types, and providers

Modules can also contain custom facts, custom functions, custom types, and providers. For more information about these, refer to [Lab 9], *[External Tools and the Puppet Ecosystem]*.

Third-party modules

You can download modules provided by other people and use them in your own manifests just like the modules you create.

Module organization

For more details on how to organize your modules, see the puppetlabs website:
https://puppet.com/docs/puppet/latest/modules_fundamentals.html.

Using standard naming conventions

Choosing appropriate and informative names for your modules and classes will be a big help when it comes to maintaining your code. This is even truer if other people need to read and work on your manifests.

How to do it...

Here are some tips on how to name things in your manifests:

1. Name modules after the software or service they manage, for example, `apache` or `haproxy`.
2. Name classes within modules (subclasses) after the function or service they provide to the module, for example, `apache::vhosts` or `rails::dependencies`.
3. If a class within a module disables the service provided by that module, name it `disabled`. For example, a class that disables Apache should be named `apache::disabled`.
4. Create a roles and profiles hierarchy of modules. Each node should have a single role consisting of one or more profiles. Each profile module should configure a single service.
5. The module that manages users should be named `user`.
6. Within the `user` module, declare your virtual users within the `user::virtual` class (for more on virtual users and other resources, see the *[Using virtual resources] recipe in [Lab 5], [Users and Virtual Resources]*).
7. Within the `user` module, subclasses for particular groups of users should be named after the group, for example, `user::sysadmins` or `user::contractors`.
8. When using Puppet to deploy the config files for different services, name the file after the service, but with a suffix indicating what kind of file it is, for example:
 - Apache `init` script: `apache.init`
 - Logrotate config snippet for Rails: `rails.logrotate`
 - Nginx `vhost` file for mywizzoapp: `mywizzoapp.vhost.nginx`
 - MySQL config for standalone server: `standalone.mysql`
 - If you need to deploy a different version of a file depending on the operating system release, for example, you can use a naming convention like the following:

```
memcached.lucid.conf
memcached.precise.conf
```

- You can have Puppet automatically select the appropriate version as follows:

```
source = > "puppet:///modules/memcached /memcached.${::lsbdistrelease}.conf",
```

9. If you need to manage, for example, different Ruby versions, name the class after the version it is responsible for; for example, `ruby192` or `ruby186`.

There's more...

Some people prefer to include multiple classes on a node by using a comma-separated list, rather than separate include statements; for example:

```
node 'server014' inherits 'server' {  
  include mail::server, repo::gem, repo::apt, zabbix  
}
```

This is a matter of style, but I prefer to use separate include statements, one on a line, because it makes it easier to copy and move around class inclusions between nodes without having to tidy up the commas and indentation every time.